

Homework 1  
CSC 277 / 477  
End-to-end Deep Learning  
Fall 2025

John Doe - `jdoe@ur.rochester.edu`

**Deadline:** See Blackboard

Instructions

Your homework solution must be typed and prepared in  $\text{\LaTeX}$ . It must be output to PDF format. To use  $\text{\LaTeX}$ , we suggest using `http://overleaf.com`, which is free.

Your submission must cite any references used (including articles, books, code, websites, and personal communications). All solutions must be written in your own words, and you must program the algorithms yourself. **If you do work with others, you must list the people you worked with.** Submit your solutions as a PDF to Blackboard.

Your programs must be written in Python. You will use *PyTorch with Lightning Fabric* (PyTorch, Lightning Fabric) in this homework. All deliverables should be within the answer box provided, i.e., `\begin{answerbox} DELIVERABLES \end{answerbox}`. If a problem requires code as a deliverable, then the code should be shown as part of the solution. One easy way to do this in  $\text{\LaTeX}$  is to use the verbatim environment, i.e., `\begin{verbatim} YOUR CODE \end{verbatim}`.

**About Homework 1:** Homework 1 aims to acquaint you with hyperparameter tuning, network fine-tuning, WandB for Training Monitoring, and model testing. *Keep in mind that network training is time-consuming, so begin early!* Copy and paste this template into an editor, e.g., `www.overleaf.com`, and then just type the answers in. You can use a math editor to make this easier, e.g., CodeCogs Equation Editor or MathType. You may use the AI (LLM) plugin for Overleaf for help you with  $\text{\LaTeX}$ formatting.

作业1 CSC 277 /  
477 端到端深度学习  
2025年秋季

约翰·多伊 - `jdoe@ur.rochester.edu`

**截止日期:** 见黑板

说明

你的作业解决方案必须键入并使用 $\text{\LaTeX}$ 准备。它必须输出为PDF格式。为了使用 $\text{\LaTeX}$ ，我们建议使用 `http://overleaf.com`，这是免费的。

你的提交必须引用所有使用的参考文献（包括文章、书籍、代码、网站和个人通信）。所有解决方案必须用你自己的话写，并且你必须自己编写算法。**如果你与他人合作，你必须列出你合作的人。**将你的解决方案作为PDF提交到黑板。

你的程序必须用Python编写。你将在这次作业中使用*PyTorch with Lightning Fabric*(PyTorch, 闪电织物)。所有交付内容应在提供的答案框内，即`\begin{answerbox}` 交付内容 `\end{answerbox}`。如果一个问题的交付内容需要代码，那么代码应作为解决方案的一部分展示。在 $\text{\LaTeX}$ 中，一个简单的方法是使用verbatim环境，即`\begin{verbatim}` 你的代码 `\end{verbatim}`。

**关于作业1:** 作业1旨在让你熟悉超参数调优、网络微调、使用WandB进行训练监控以及模型测试。请记住，网络训练非常耗时，所以请尽早开始！将此模板复制粘贴到编辑器中，例如 `www.overleaf.com`，然后只需输入答案即可。你可以使用数学编辑器来简化这个过程，例如CodeCogs公式编辑器或MathType。你可以使用AI（大型语言模型）插件帮助你在Overleaf中进行 $\text{\LaTeX}$ 格式化。

## Problem 1 - WandB for Training Monitoring

Training neural networks involves exploring different model architectures, hyperparameters, and optimization strategies. Monitoring these choices is crucial for understanding and improving model performance. Logging experiment results during training helps to:

- Gain insights into model behavior (e.g., loss, accuracy, convergence patterns).
- Optimize hyperparameters by evaluating their impact on stability and accuracy.
- Detect overfitting or underfitting and make necessary adjustments.

In this problem, you'll train ResNet-18 models for image classification on the Oxford-IIIT Pet Dataset while exploring various hyperparameters. You'll use Weights and Biases (W&B) to log your experiments and refine your approach based on the results.

### Part 1: Implementing Experiment Logging with W&B (7 points)

**Prepare the Dataset.** Download the dataset. Complete the dataset definition in `train.py`. Make a custom `Dataset` class (examples, reference) or use `ImageFolder` to do this. Use the splits defined in `oxford_pet_split.csv` to define the training, validation, and test splits. In the dataset's preprocessing, resize the images to 224 as required by ResNet-18, and apply image normalization using statistics from the training set or from ImageNet.

**Evaluating Model Performance.** During model training, the validation set is a crucial tool to prevent overfitting. Complete `evaluate()` function in `train.py` which takes a model and a dataloader as inputs and outputs the model's accuracy score and cross-entropy loss on the dataset.

**Integrate W&B Logging.** To integrate W&B for experiment logging, follow these steps and add the necessary code to `train.py`:

1. Refer to the W&B official tutorial for guidance.
2. Initialize a new run at the start of the experiment following the tutorial's code snippet. Log basic experiment **configurations**, such as total training epochs, learning rate, batch size, and scheduler usage. Ensure the run **name** is interpretable and reflects these key details.
3. During training, log the training loss and learning rate after each mini-batch.
4. After each epoch, log the validation loss and validation accuracy.
5. At the end of the training, log the model's performance on the test set, including loss and accuracy scores.

**Experiment and Analysis.** Execute the experiment using the **default** setup (found in `get_args` function). Log in to the W&B website to inspect your implementation.

#### Deliverable:

- Code snippet(s) completed to load the dataset.
- Screenshot(s) of the experiment configuration (*Select the specific Run in W&B, and then the Overview tab*)

## 问题 1 - 使用WandB进行训练监控

训练神经网络涉及探索不同的模型架构、超参数和优化策略。监控这些选择对于理解和改进模型性能至关重要。在训练过程中记录实验结果有助于：

- 深入了解模型行为（例如，损失、准确率、收敛模式）。
- 通过评估超参数对稳定性和准确率的影响来优化超参数。
- 检测过拟合或欠拟合并进行必要的调整。

在这个问题中，你将使用ResNet-18模型对牛津-IIIT宠物数据集进行图像分类训练，同时探索各种超参数。你将使用权重和偏差（W&B）来记录你的实验，并根据结果优化你的方法。

### P第1部分：使用W&B实现实验记录（7分）

**准备数据集。** 下载数据集。在 `train.py` 中完成数据集定义。创建一个自定义的 `Dataset` 类（示例，参考）或使用 `ImageFolder` 来完成此操作。使用 `oxford_pet_split.csv` 中定义的分割来定义训练、验证和测试分割。在数据集的预处理中，将图像调整到 ResNet-18 所需的 224 大小，并使用训练集或 ImageNet 的统计数据应用图像归一化。

**评估模型性能。** 在模型训练过程中，验证集是防止过拟合的关键工具。在 `train.py` 中完成 `evaluate()` 函数，该函数接受一个模型和一个数据加载器作为输入，并输出模型在数据集上的准确度分数和交叉熵损失。

**集成W&B记录。** 为了集成W&B进行实验记录，按照以下步骤操作，并在 `train.py` 中添加必要的代码：

1. 参考W&B官方教程以获取指导。
2. 在实验开始时按照教程的代码片段初始化一个新的运行。记录基本的实验**配置**，例如总训练周期、学习率、批量大小和调度器使用。确保运行**名称**是可解释的，并反映这些关键细节。
3. 在训练过程中，在每个小批量后记录训练损失和学习率。
4. 在每个周期结束后，记录验证损失和验证准确度。
5. 在训练结束时，记录模型在测试集上的性能，包括损失和准确度分数。

**实验与分析。** 使用默认设置（在 `get_args` 函数中找到）执行实验。登录W&B网站以检查您的实现。

#### 交付内容：

- 完成加载数据集的代码片段。
- 实验配置的截图（在W&B中选择特定的运行，然后选择概览标签页）

- Screenshot(s) of all logged charts (*Select the specific Run in W&B, and then Charts under Workspace tab*).
- Are the data logged accurately in the W&B interface? Does the experiment configuration align with your expectations?
- Analyze the logged charts to determine whether the training has converged.

Answer:

Part 2: Tuning Hyperparameters

In this section, you’ll experiment with key hyperparameters like learning rate and scheduler, and study their effects on training dynamics. For each step, change only one configuration at a time. Try not modify other hyperparameters (except batch size, which can be adjusted based on your computing resources).

Learning Rate Tuning with Sweep (5 points)

The learning rate is a crucial hyperparameter that significantly affects model convergence and performance. Run the training script using W&B sweep with the following learning rates:  $1e-2$ ,  $1e-4$ , and  $1e-5$ . Also, include the default learning rate ( $1e-3$ ) from Part 1 in your analysis.

Deliverable:

- Provide screenshots of logged charts showing learning rate, training loss, validation accuracy, and final test accuracy. Each chart should display results from **multiple runs** (all four learning rates in one chart). Ensure that titles and legends are clear and easy to interpret.
- Analyze how the learning rate impacts the training process and final performance.
- Code of your sweep configuration that defines the search space.

Answer:

Gradient Norm Monitoring for Exploding/Vanishing Gradients (4 points)

Monitoring gradient norms during training is essential to detect exploding gradients (leading to instability or NaN losses) or vanishing gradients (causing slow/no convergence), which are common critical issues in deep learning practice. In `train.py`, extend your W&B logging to compute and log the average L2 gradient norm across all model parameters after each batch (e.g., via a loop over `model.parameters()`). Run a ‘problematic’ variant (e.g., set learning rate to 0.1 or larger values to induce exploding gradients). Compare the norm charts and loss curves of the problematic variant and your default run.

- 所有记录的图表的截图（在W&B中选择特定的运行，然后在工作区标签下选择图表）。
- 数据是否在W&B界面中准确记录？实验配置是否符合您的预期？
- 分析记录的图表，以确定训练是否已收敛。

答案:

第二部分：调整超参数

在本节中，您将实验关键超参数，如学习率和调度器，并研究它们对训练动态的影响。对于每个步骤，一次只更改一个配置。尽量不要修改其他超参数（批量大小除外，可以根据您的计算资源进行调整）。

使用Sweep进行学习率调优（5分）

学习率是一个关键的超参数，它显著影响模型的收敛和性能。使用W&B sweep运行训练脚本，并采用以下学习率： $1e-2$ 、 $1e-4$ 和 $1e-5$ 。同时，在您的分析中包含第一部分的默认学习率（ $1e-3$ ）。

交付内容:

- 提供记录的图表截图，显示学习率、训练损失、验证准确度和最终测试准确度。每个图表应展示**multiple runs**（所有四个学习率在一个图表中）的结果。确保标题和图例清晰且易于解读。
- 分析学习率如何影响训练过程和最终性能。
- 定义搜索空间的扫描配置代码。

答案:

梯度范数监控用于爆炸/消失梯度（4分）

在训练过程中监控梯度范数对于检测爆炸梯度（导致不稳定或NaN损失）或消失梯度（导致收敛缓慢/不收敛）至关重要，这些是深度学习实践中常见的严重问题。在 `train.py`中，扩展您的W&B记录以计算并在每个批次后记录所有模型参数的平均L2梯度范数（例如，通过循环遍历 `model.parameters()`）。运行一个‘问题’变体（例如，将学习率设置为0.1或更大值以引发爆炸梯度）。比较问题变体和默认运行的范数图表和损失曲线。

**Deliverable:**

- Screenshots of W&B charts showing gradient norms and training loss over steps for both runs.
- Do you notice signs of exploding/vanishing gradients in the problematic run? If so, identify them and explain your inferences.
- Code snippet(s) to compute and log gradient norms.

**Hint:** Consider carefully the placement of the Gradient norm calculation in the training loop’s backward pass.

**Answer:**

**Learning Rate Scheduler (4 points)**

Learning rate schedulers dynamically adjust the learning rate during training, improving efficiency, convergence, and overall performance. In this step, you’ll implement the `OneCycleLR` scheduler in the `get_scheduler()` function within `train.py`. Compare the results to the baseline (default setting). If implemented correctly, the learning rate will initially increase and then decrease during training.

**Deliverable:**

- Provide charts comparing the new setup with the baseline: learning rate, training loss, validation accuracy, and final test accuracy.
- Explain how the `OneCycleLR` scheduler impacts the learning rate, training process, and final performance compared to the baseline.

**Answer:**

**Part 3: Scaling Learning Rate with Batch Size (4 points)**

As observed in previous parts, the choice of learning rate is crucial for effective training. As batch size increases, the effective step size in the parameter space also increases, requiring adjustments to the learning rate. In this section, you’ll investigate how to scale the learning rate appropriately when the batch size changes. Read the first few paragraphs of this blog post to understand scaling rules for Adam (used in default) and SGD optimizers. Then, conduct experiments to verify these rules. First, double (or halve) the batch size without changing the learning rate and run the training script. Next, ONLY adjust the learning rate as suggested in the post. Compare these results with the default setting. Note that since the total training steps vary with batch size, you should also log the number of seen examples to create accurate charts for comparison.

**Deliverable:**

**交付内容:**

- 显示两个运行过程中梯度范数和训练损失随步骤变化的W&B图表的截图。
- 你在问题运行中注意到爆炸/消失梯度的迹象了吗？如果有，请识别它们并解释你的推断。
- 计算和记录梯度范数的代码片段。

**提示:** 仔细考虑梯度范数计算在训练循环反向传播中的位置。

**答案:**

**学习率调度器 (4分)**

学习率调度器在训练过程中动态调整学习率，提高效率、收敛性和整体性能。在这一步中，你将在 `train.py`中的 `get_scheduler()` 函数中实现`OneCycleLR` 调度器。将结果与基线（默认设置）进行比较。如果正确实现，学习率将在训练初期增加，然后逐渐减少。

**交付内容:**

- 提供图表，比较新设置与基线：学习率、训练损失、验证准确度和最终测试准确度。
- 解释 `OneCycleLR` 调度器如何影响学习率、训练过程和最终性能，与基线相比。

**答案:**

**第三部分：随批量大小缩放学习率（4分）**

如前几部分所观察到的，学习率的选择对于有效训练至关重要。随着批量大小的增加，参数空间中的有效步长也会增加，这需要调整学习率。在本节中，你将研究在批量大小变化时如何适当地缩放学习率。阅读这篇博客文章的前几段，以了解Adam（默认使用）和SGD优化器的缩放规则。然后，进行实验以验证这些规则。首先，将批量大小加倍（或减半），但不改变学习率，并运行训练脚本。接下来，仅根据文章中的建议调整学习率。将这些结果与默认设置进行比较。请注意，由于总训练步数随批量大小而变化，因此你还应记录已见示例的数量，以创建用于比较的准确图表。

**交付内容:**

- Present charts showing: training loss and validation accuracy (with the x-axis being `seen_examples`), and final test accuracy. Ensure the legends are clear. You may apply smoothing for better visualization.
- Analyze the results: do they align with the patterns discussed in the blog post?

Answer:

Part 4: Fine-Tuning a Pretrained Model (3 points)

Fine-tuning leverages the knowledge of models trained on large datasets by adapting their weights to a new task. In this section, you will fine-tune a ResNet-18 model pre-trained on ImageNet using `torchvision.models.resnet18(pretrained=True)`. Modify the classification head to match the number of classes in your task, and replace the model definition in the original code. Keep the rest of the setup as default for comparison.

Deliverable:

- Present charts showing: training loss, validation accuracy, and final test accuracy.
- Analyze the impact of pre-training on the model’s learning process and performance.

Answer:

- 展示图表，显示：训练损失和验证准确度（x轴为`seen_examples`），以及最终测试准确度。确保图例清晰。你可以应用平滑以获得更好的可视化效果。
- 分析结果：它们是否符合博客文章中讨论的模式？

答案:

第4部分：微调预训练模型（3分）

微调利用在大数据集上训练的模型的知识，通过调整其权重以适应新任务。在本节中，你将使用 `torchvision.models.resnet18(pretrained=True)` 微调一个在ImageNet上预训练的ResNet-18模型。修改分类头以匹配你任务中的类别数量，并替换原始代码中的模型定义。为了比较，保持其余设置默认。

交付内容:

- 展示图表，显示：训练损失、验证准确度和最终测试准确度。
- 分析预训练对模型学习过程和性能的影响。

答案:



## Problem 2 - Model Testing

Unlike model evaluation, which focuses on performance metrics, model testing ensures that a model behaves as expected under specific conditions.

- **Pre-Train Test:** Conducted before training, these tests identify potential issues in the model’s architecture, data preprocessing, or other components, preventing wasted resources on flawed training.
- **Post-Train Test:** Performed after training, these tests evaluate the model’s behavior across various scenarios to ensure it generalizes well and performs as expected in real-world situations.

In this problem, you will examine the code and model left by a former employee who displayed a lack of responsibility in his work. The code can be found in the `Problem 2` folder. The necessary predefined functions for this task are available in the `model_testing.py` file. Follow the instructions provided in that file for detailed guidance.

### Part 1: Pre-Train Testing

In this part, you will apply pre-train tests to verify experimental correctness before training is conducted. **Deliverables:** For each question in Part 1, provide clear deliverables of the following:

1. Observations and analysis of the results.
2. Suggested approaches for addressing the detected issues (if any).
3. Code implementation.

### Data Leakage Check (3 points)

Load the training, validation, and test data sets using `get_dataset()` function. Check for potential data leakage between these sets by directly comparing the images, as data augmentation was not applied. Since identical objects usually have different hash values in Python, consider using techniques like image hashing for this comparison.

Answer:

### Model Architecture Check (2 points)

Initialize the model using the `get_model()` function. Verify that the model’s output shape matches the label format (hint: consider the number of classes in the dataset).

Answer:

## 问题 2 - 模型测试

与关注性能指标的模型评估不同，模型测试确保模型在特定条件下按预期行为。

- **预训练测试:** 在训练之前进行，这些测试识别模型架构、数据预处理或其他组件中的潜在问题，防止在存在缺陷的训练上浪费资源。
- **后训练测试:** 在训练之后进行，这些测试评估模型在各种场景下的行为，以确保其泛化良好，并在真实世界情境中按预期表现。

在这个问题中，你将检查一位前员工留下的代码和模型，该员工在工作中表现出缺乏责任感。代码可以在 `Problem 2` 文件夹中找到。此任务所需的预定义函数位于 `model_testing.py` 文件中。请按照该文件中提供的说明进行详细指导。

### 第一部分：预训练测试

在这一部分，你将应用预训练测试来验证实验正确性，在训练进行之前。**交付物：**对于第一部分的每个问题，提供以下清晰的交付物：

1. 结果的观察和分析。
2. 针对检测到的问题（如果有）的 **suggested approaches**。
3. 代码实现。

### 数据泄露检查（3分）

使用 `get_dataset()` 函数加载训练、验证和测试数据集。通过直接比较图像来检查这些集合之间的潜在数据泄露，因为未应用数据增强。由于在Python中相同的对象通常具有不同的哈希值，考虑使用图像哈希等技术进行此比较。

答案：

### 模型架构检查（2分）

使用 `get_model()` 函数初始化模型。验证模型的输出形状是否与标签格式匹配（提示：考虑数据集中的类别数量）。

答案：

Gradient Descent Validation (2 points)

Verify that ALL the model’s trainable parameters are updated after a single gradient step on a batch of data.

Answer:

Learning Rate Check (2 points)

Implement the learning rate range test using pytorch-lr-finder. Determine whether the learning rate is appropriately set by examining the loss-learning rate graph. Necessary components for `torch_lr_finder.LRFinder` are provided in `model_testing.py`.

Answer:

Part 2: Post-Train Testing

Dying ReLU Examination (4 points)

In this section, you will examine the trained model for “Dying ReLU.” Dying ReLU occurs when a ReLU neuron outputs zero consistently and cannot differentiate between inputs. Load the trained model using `get_trained_model()` function, and the test set using `get_test_set()` function. Review the model’s architecture, which is based on ResNet and can be found in `utils/trained_models.py`. Then address the following:

- 1. Identify the layer(s) where Dying ReLU might occur and explain why.
- 2. Describe your approach for detecting Dying ReLU neurons.
- 3. Determine if Dying ReLU neurons are present in the trained model, and provide your code implementation..

Hint: Consider how BatchNorm operation would influence the presence of dying ReLU.

Answer:

Model Robustness Test - Brightness (4 points)

In this section, you will evaluate the model’s robustness to changes in image brightness using a defined brightness factor. Define a brightness factor  $\lambda$ , which determines the image brightness by multiplying pixel values by  $\lambda$ . Specifically,  $\lambda = 1$  corresponds to the original image’s brightness. Load the trained model using `get_trained_model()` function, and the test dataset using `get_test_set()` function. Investigate the model’s performance across various brightness levels by adjusting  $\lambda$  from 0.2 to 1.0 in increments of 0.2.

梯度下降验证 (2分)

验证在单个数据批次的梯度步长后，模型的所有可训练参数是否都已更新。

答案：

学习率检查 (2分)

使用pytorch-lr-finder实现学习率范围测试。通过检查损失-学习率图来确定学习率是否设置得当。 `torch_lr_finder.LRFinder` 所需的组件在 `model_testing.py`中提供。

答案：

第2部分：训练后测试

死亡ReLU检查（4分）

在本节中，你将检查训练模型是否存在“死亡ReLU”。死亡ReLU是指ReLU神经元持续输出零，无法区分输入的情况。使用 `get_trained_model()` 函数加载训练模型，使用 `get_test_set()` 函数加载测试集。审查基于ResNet的模型架构，该架构可以在 `utils/trained_models.py`中找到。然后回答以下问题：

- 1. 识别可能发生死亡ReLU的层，并解释原因。
- 2. 描述你检测死亡ReLU神经元的 **approach**。
- 3. 确定训练模型中是否存在死亡ReLU神经元，并提供你的代码实现。

提示：考虑批量归一化操作如何影响死亡ReLU的出现。

答案：

模型鲁棒性测试 - 亮度（4分）

在本节中，你将使用定义的亮度因子评估模型对图像亮度变化的鲁棒性。定义一个亮度因子  $\lambda$ ，通过将像素值乘以  $\lambda$  来确定图像亮度。具体来说， $\lambda = 1$  对应于原始图像的亮度。使用 `get_trained_model()` 函数加载训练模型，并使用 `get_test_set()` 函数加载测试数据集。通过将  $\lambda$  从 0.2 调整到 1.0，以 0.2 的增量来研究模型在不同亮度级别上的性能。

**Deliverable:**

- 1. Plot a curve showing how model accuracy varies with brightness levels.
- 2. Analyze the relationship and discuss any trends observed.

**Answer:**

**Model Robustness Test - Rotation (4 points)**

Evaluate the model’s robustness to changes in image rotation. Rotate the input image from 0 to 300 degrees in increments of 60 degrees. Similarly, load the trained model using `get_trained_model()` function, and the test set using `get_test_set()` function.

**Deliverable:**

- 1. Plot a curve showing the relationship between rotation angles and model accuracy.
- 2. Analyze the trend and discuss any observed patterns.
- 3. Suggest potential improvements to enhance model robustness

**Answer:**

**Normalization Mismatch (2 points)**

Load the test set using the `get_test_set()` function. Assume that the mean and standard deviation (std) used to normalize the testing data are different from those applied to the training data.

**Deliverable:**

- 1. Calculate and report the mean and std of the images in the loaded test set (tutorial). Compare these values with the expected mean and std after proper normalization.
- 2. Discuss one potential impact of this incorrect normalization on the model’s performance or predictions.

**Answer:**

**交付内容:**

- 1. 绘制一条曲线，显示模型准确性如何随亮度级别变化。
- 2. 分析关系并讨论观察到的任何趋势。

**答案:**

**模型鲁棒性测试 - 旋转（4分）**

评估模型对图像旋转变化的鲁棒性。将输入图像从0度旋转到300度，每次增加60度。同样，使用`get_trained_model()` 函数加载训练模型，使用 `get_test_set()` 函数加载测试集。

**交付内容:**

- 1. 绘制一条曲线，显示旋转角度与模型准确性之间的关系。
- 2. 分析趋势并讨论任何观察到的模式。
- 3. 提出潜在改进以增强模型鲁棒性

**答案:**

**归一化不匹配（2分）**

使用 `get_test_set()` 函数加载测试集。假设用于归一化测试数据的均值和标准差（std）与应用于训练数据的均值和标准差不同。

**交付内容:**

- 1. 计算并报告加载的测试集（教程）中图像的均值和标准差。将这些值与适当归一化后的预期均值和标准差进行比较。
- 2. 讨论这种不正确归一化对模型性能或预测的一个潜在影响。

**答案:**